

PDF-Based Question Answering Chatbot

Project Overview

This project is a PDF-based Question Answering (QA) chatbot that allows users to ask questions based on the content of a provided PDF document. If the chatbot cannot find a relevant answer, it gracefully handles the situation with a fallback response.

Approach

The development of the PDF-based Question Answering Chatbot was guided by the following systematic approach:

1. Problem Understanding

- Clearly defined the goal: A chatbot capable of answering questions based on a provided PDF document.
- Identified two key challenges:
 - Extracting and organizing content from the PDF.
 - Handling queries with appropriate fallback responses.
 - Building backend and frontend for making a complete system

2. Tool and Library Selection

- PDF Parsing: Selected PDFMiner for extracting text from PDF documents.
- Embedding and Search:
 - HuggingFace Bge Embeddings: For generating semantic embeddings of the document content.
 - ChromaDB: For efficient storage and similarity-based retrieval of document chunks.
- Chat Interface: Chose a simple JavaScript-based frontend for user interaction.

- Backend Framework: Used FastAPI to implement the chatbot's API.
 - AI Prompting: Utilized LangChain for managing AI prompt templates and generating responses.
-

3. Document Preparation

- PDF Loading: Loaded the content of the PDF using PyPDFMinerLoader.
 - Chunking: Split the document into semantically similar chunks to ensure better query matching by using semantic-text-splitter, using a token limit for efficient processing. Chose this over traditional splitting as this generally gives better retrieved documents
 - Embedding: Converted chunks into semantic vectors using HuggingFace Bge Embeddings for indexing in ChromaDB.
-

4. Semantic Search Implementation

- Configured ChromaDB to retrieve the most relevant document chunks for a given query.
- Set a **similarity score threshold** to filter out irrelevant matches, ensuring only meaningful context is considered. This is done in order to ensure that application does not return any irrelevant responses and gives a fallback answer if the model cannot find the answer in the document
- Fallback returned is

“Sorry, I didn’t understand your question. Do you want to connect with a live agent?”

5. Chatbot Response Generation

- Designed a prompt template in LangChain that:
 - Provides the retrieved context for the AI to answer the query.
 - Implements fallback logic if the retrieved context is irrelevant.
 - The prompt is designed in such a way that the model only returns answer if it is confident about it , otherwise a fallback response is given
- Integrated the prompt template with AI model to dynamically generate responses.

- Groq is used to get response from the AI model through its API
 - *Mixtral-8x7b-32768* model is used because of its great accuracy in question answering tasks
-

6. Frontend Development

- Built a simple web interface using JavaScript, allowing users to:
 - Input their questions.
 - View the chatbot's responses.
 - Handle loading and error states.
-

7. Backend-Frontend Integration

- Connected the frontend to the backend via API endpoints hosted on FastAPI.
 - Configured the frontend to make POST requests to the backend with the conversation data and display the responses.
-

8. Fallback Logic

- Implemented fallback logic to handle cases where no relevant context is retrieved from the document:
 - Checked the similarity score of retrieved chunks.
 - Responded with:
"Sorry, I didn't understand your question. Do you want to connect with a live agent?"
 - Ensured this logic was consistent in both the backend API and the frontend UI.
 - Also added a safety check in *prompt* so that the model returns fallback if the context is not present in the pdf
-

9. Testing and Debugging

- Conducted extensive testing for:
 - PDF parsing and chunking accuracy.
 - Query matching and retrieval quality.

- AI-generated responses.
 - Seamless frontend-backend communication.
 - Resolved edge cases such as invalid queries, empty PDF content, and connection errors.
 - The *score threshold* is set to **0.5** through testing with different human made queries and found this threshold to perform decent
-

10. Deployment

- Developed a Python script to automate the starting of both backend and frontend servers.
 - Ensured the application could be run locally with minimal setup.
-

Technologies Used

Frontend

- HTML, CSS, JavaScript: For the web interface.
- HTTP Server: Serves the static frontend files.

Backend

- FastAPI: For building the chatbot API.
- LangChain: For handling document splitting, vector embeddings, and semantic search.
- ChromaDB: For storing and querying the document embeddings.
- HuggingFace Embeddings: For semantic vector generation.
- PDFMiner: For parsing and extracting text from PDF documents.

Other Tools

- Uvicorn: For running the FastAPI application.
 - Python HTTP Server: For hosting the frontend.
-

Installation and Setup

Prerequisites

- Python 3.7 or higher
- pip (Python package manager)

Project structure

project_directory/

```
|---- backend/  
|-----__init__.py  
|----- app.py  
|-----input_apport.pdf
```

```
|---- frontend/  
|-----index.html  
|-----app.js  
|-----style.css  
|-----assets/
```

```
|— startup.py  
|-----.env  
|-----.gitignore  
|-----requirements.txt
```

Follow these steps to set up and run the PDF-based Question Answering Chatbot:

1. Open the Project in Visual Studio Code

1. Open Visual Studio Code on your computer.
 2. Use the menu option `File > Open Folder` (or `Open on macOS`).
 3. Navigate to the project directory and select it.
-

2. Create and Activate a Virtual Environment

For macOS/Linux:

1. Open the terminal in Visual Studio Code or your system terminal.
2. Run the following commands to create and activate the virtual environment:

```
python3 -m venv venv  
source venv/bin/activate
```

For Windows:

1. Open the terminal in Visual Studio Code or the Command Prompt.
2. Run the following commands to create and activate the virtual environment:

```
python -m venv venv  
venv\Scripts\activate
```

3. Install Project Dependencies

1. Ensure the virtual environment is activated.
2. Install the dependencies specified in `requirements.txt` by running:

```
pip install -r requirements.txt
```

3. Verify that all dependencies have been installed without errors.
4. Run the application using the following command

```
python start_app.py
```

5. Access the Application:

- Backend API Documentation: <http://127.0.0.1:8000/docs>
 - Frontend Interface: <http://127.0.0.1:5500/>
-

Usage Instructions

1. Upload or specify the PDF document (`input_apport.pdf`) before starting the server.
 2. Open the frontend interface in your browser.
 3. Type a question related to the content of the PDF.
 4. Receive a response:
 - If the answer is found, it is displayed.
 - If no answer is found, you will see the fallback message:
"Sorry, I didn't understand your question. Do you want to connect with a live agent?"
-

Troubleshooting

If the application is not starting using `startup.py`, manually start backend and frontend using the below steps:

Start the Backend

1. While the virtual environment is activated, navigate to the project directory if not already there.
2. Run the following command to start the backend server:

```
uvicorn backend.app:app --reload --host 127.0.0.1 --port 8000
```

The backend server should now be running on <http://127.0.0.1:8000>

. Start the Frontend

1. Open another terminal instance (or tab) while keeping the backend terminal running.
2. Navigate to the frontend directory within the project

```
cd frontend
```

3. Start a simple HTTP server to serve the frontend files:

```
python -m http.server 5500
```

4. The frontend should now be running on <http://127.0.0.1:5500>

Acknowledgments

- [LangChain](#) for document processing and querying.
 - [HuggingFace](#) for embedding models.
 - FastAPI for building the API.
-